




Software Solutions to Micro-architectural Side Channels

Yinqian Zhang
Assistant Professor
Computer Science & Engineering
The Ohio State University

Introduction

- Research interests
 - Computer system security in general
 - (Micro-architectural) side-channel attacks and defenses
- Recent publications on micro-architectural side channels
 - Cloud computing (S&P'11, CCS'12, CCS'13, CCS'14, Security'15, Security'16, RAID'16, CCS'16a, AsiaCCS'17b, TDSC 2018)
 - Intel SGX (AsiaCCS'17a, CCS'17a, CCS'17b, S&P'18a, SgxPectre)
 - Smartphones (CCS'16b), side-channel measurement (ACSAC'18)
- Recent publications on system-level side channels
 - Mobile systems (CCS'15, NDSS'18a), searchable encryption (INFOCOM'18), vulnerability detection (S&P'18b)
- Service on the program committees of security conferences
 - IEEE S&P (2016, 2017, 2018), ACM CCS (2015, 2016, 2017, 2018), Usenix Security (2017), NDSS (2017, 2018, 2019)

Hyper-Threading Side Channels



Physical CPU Core


Hyper-Threading creates side channels

- Floating point units
- Branch prediction units
- Store buffers
- Translation lookaside buffers
- Caches (L1I, L1D, L2)

Cache missing for fun and profit, Colin Percival 2005

Cache Attacks and Countermeasures: the Case of AES, Dag Arne Osvik, Adi Shamir, and Eran Tromer, 2005

Hyper-Threading Side Channels



Hyper-Threading enhances existing side channels

- Page monitoring through page tables (Wang et al. 2017)


Hyper-Threading thwarted defenses

- Hyper-Threading invalidates some defense techniques that leverage Intel's TSX (Gruss et al. 2017)

Leaky Cauldron on the Dark Land: Understanding Memory Side-Channel Hazards in SGX, Wenhao Wang, Guoxing Chen, Xiaorui Pan, Yinqian Zhang, XiaoFeng Wang, Vincent Bindschaedler, Haixu Tang, Carl A. Gunter, *ACM Conference on Computer and Communications Security*, Dallas, Texas, USA, Oct. 2017

Strong and Efficient Cache Side-Channel Protection using Hardware Transactional Memory, Daniel Gruss, Julian Lettner, Felix Schuster, Olya Ohrimenko, Istvan Haller, Manuel Costa, *USENIX Security Symposium*, 2017.

Hyper-Threading Side Channels



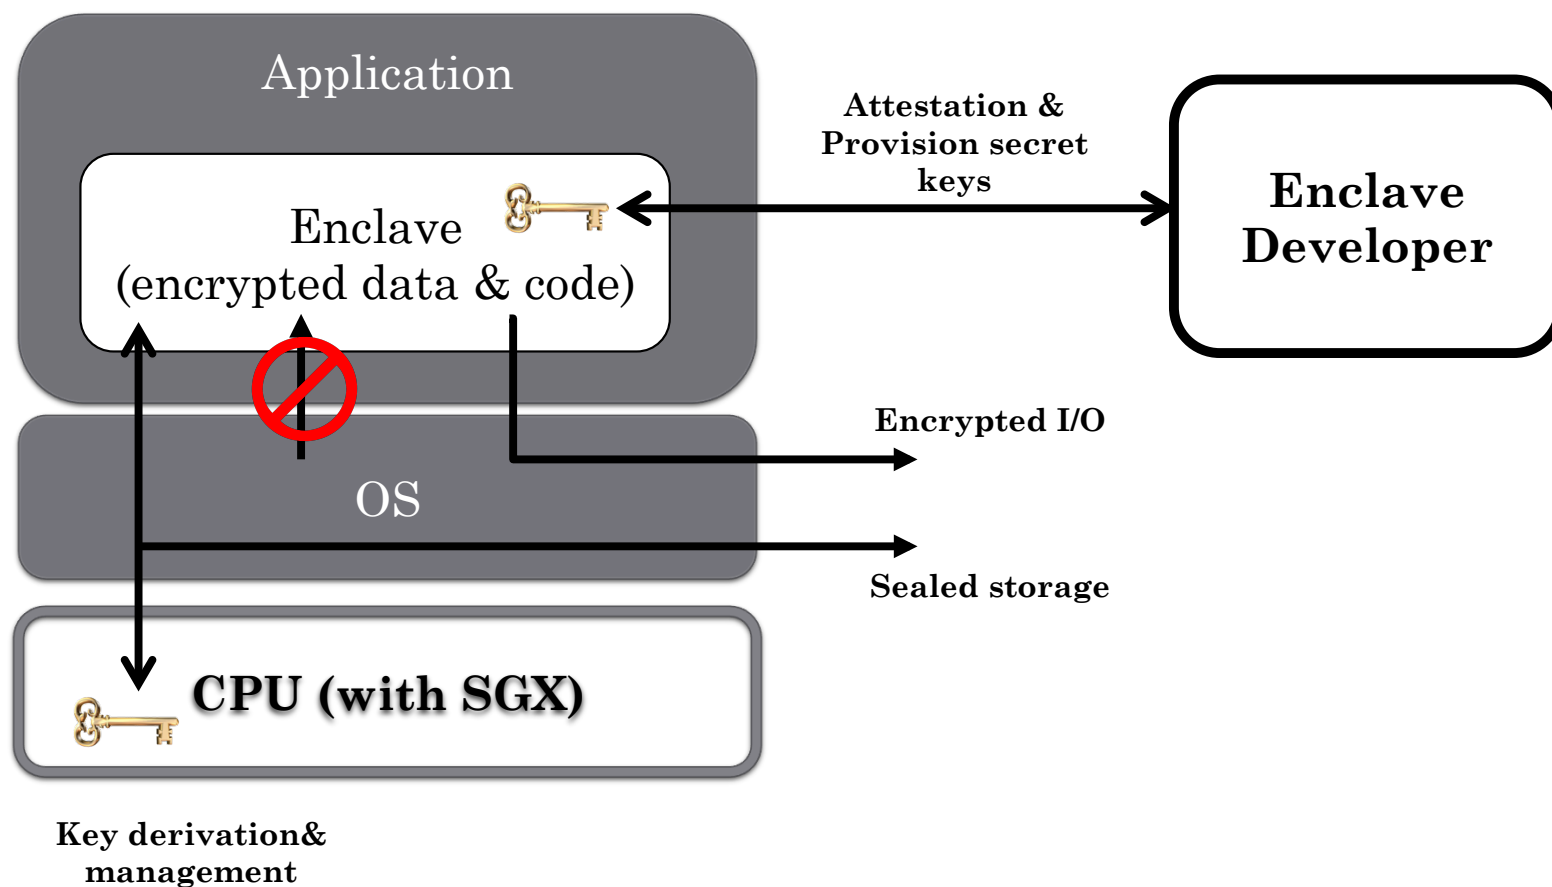
Hyper-Threading facilities speculative-execution based side channels

- Meltdown attacks (Lipp et al. 2018)
- Foreshadow attacks / L1 Terminal Fault (Bulck et al. 2018)

Meltdown: Reading Kernel Memory from User Space, Moritz Lipp and Michael Schwarz and Daniel Gruss and Thomas Prescher and Werner Haas and Anders Fogh and Jann Horn and Stefan Mangard and Paul Kocher and Daniel Genkin and Yuval Yarom and Mike Hamburg, USENIX Security Symposium, 2018.

Foreshadow: Extracting the Keys to the Intel SGX Kingdom with Transient Out-of-Order Execution, Van Bulck, Jo and Minkin, Marina and Weisse, Ofir and Genkin, Daniel and Kasikci, Baris and Piessens, Frank and Silberstein, Mark and Wenisch, Thomas F. and Yarom, Yuval and Strackx, Raoul, USENIX Security Symposium, 2018.

Intel Software Guard Extension (SGX)



Hyper-Threading side channels on SGX

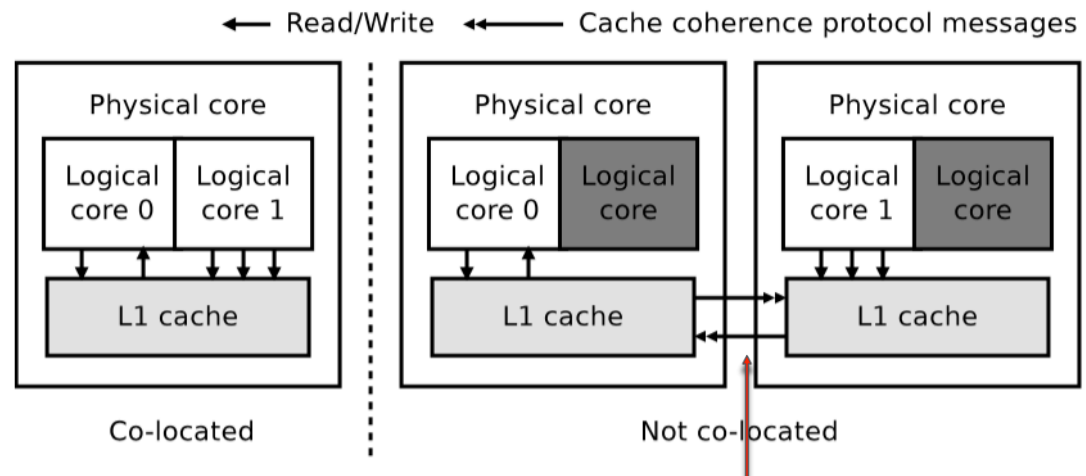
- Simply disabling Hyper-Threading is not a solution
 - No effective way to verify
- **HyperRace (Chen et al. 2018)**
 - Create an auxiliary thread, called shadow thread, to occupy the sibling Hyper-Thread
 - Request the untrusted OS to schedule the protected enclave thread and the shadow enclave thread on the same physical core
 - Verify these two threads are co-located on the same physical core

Racing in Hyperspace: Closing Hyper-Threading Side Channels on SGX with Contrived Data Races, Guoxing Chen, Wenhao Wang, Tianyu Chen, Sanchuan Chen, Yinqian Zhang, XiaoFeng Wang, Ten-Hwang Lai, Dongdai Lin, *IEEE Symposium on Security and Privacy*, 2018

Hyper-Thread co-location tests

- Two threads operate on the same shared variable ν inside the enclave:


- Thread T_0
 - While (1)
 - Writes 0 to ν
 - Waits for 10 cycles
 - Reads ν
- Thread T_1
 - While (1)
 - Writes 1 to ν



Cache coherence protocol (latency ~ 190 cycles)

- Thread T_0 reads $\nu = 1$ with very high probability when co-located
- Thread T_0 reads $\nu = 0$ with very high probability when not co-located

A refined design



Thread T_0		Thread T_1	
1 <initialization>:	31 cmovl %rbx, %r10	1 <initialization>:	31 cmp \$1, %r9
2 mov \$colocation_count, %rdx	32 sub %rax, %r9	2 mov \$colocation_count, %rdx	32 ; continuous number?
3 xor %rcx, %rcx	33 cmp \$1, %r9	3 xor %rcx, %rcx	33 cmova %r11, %r10
4 ; co-location test counter	34 ; continuous number?	4 ; co-location test counter	34 add %r10, %rcx
5 <synchronization>:	35 cmova %r11, %r10	5 <synchronization>:	35 shl \$b_count, %rbx
6 ... ; acquire lock 0	36 add %r10, %rcx	6 ... ; release lock 0	36 ; bit length of \$count
7 .sync0:	37 shl \$b_count, %rbx	7 .sync2:	37 mov %rax, %r9
8 mov %rdx, (sync_addr1)	38 ; bit length of \$count	8 mov %rdx, (sync_addr0)	38 ; record the last number
9 cmp %rdx, (sync_addr0)	39 mov %rax, %r9	9 cmp %rdx, (sync_addr1)	39 <store>:
10 je .sync1	40 ; record the last number	10 je .sync3	40 mov %rsi, (%r8)
11 jmp .sync0	41 <padding instructions 0>:	11 jmp .sync2	41 <padding instructions 1>:
12 .sync1:	42 nop	12 .sync3:	42 mov (%r8), %rax
13 mfence	43 nop	13 mfence	43 lfence
14 mov \$0, (sync_addr0)	44 :	14 mov \$0, (sync_addr1)	44 mov (%r8), %rax
15 <initialize a round>:	45 nop	15 <initialize a round>:	45 lfence
16 mov \$begin0, %rsi	46 mov (%r8), %rax	16 mov \$begin1, %rsi	46 mov (%r8), %rax
17 mov \$1, %rbx	47 mov (%r8), %rax	17 mov \$1, %rbx	47 lfence
18 mfence	48 :	18 mfence	48 mov (%r8), %rax
19 mov \$addr_v, %r8	49 mov (%r8), %rax	19 mov \$addr_v, %r8	49 lfence
20 <co-location test>:	50 dec %rsi	20 <co-location test>:	50 mov (%r8), %rax
21 .L0:	51 cmp \$end0, %rsi	21 .L2:	51 lfence
22 <load>:	52 jne .L0	22 <load>:	52 dec %rsi
23 mov (%r8), %rax	53 ; finish 1 co-location test	23 mov (%r8), %rax	53 cmp \$end1, %rsi
24 <store>:	54 <all rounds finished?>:	24 <update counter>:	54 jne .L2
25 mov %rsi, (%r8)	55 ... ; release lock 1	25 mov \$0, %r10	55 ; finish 1 co-location test
26 <update counter>:	56 dec %rdx	26 mov \$0, %r11	56 <all rounds finished?>:
27 mov \$0, %r10	57 cmp \$0, %rdx	27 cmp \$end0, %rax	57 ... ; acquire lock 1
28 mov \$0, %r11	58 jne .sync0	28 ; a data race happens?	58 dec %rdx
29 cmp \$end0, %rax		29 cmovg %rbx, %r10	59 cmp \$0, %rdx
30 ; a data race happens?		30 sub %rax, %r9	60 jne .sync2



More details behind the scene

- Co-location test via statistical hypothesis testing:
 - Consider n data race unit tests: $X_j, j = 1, 2, \dots, n$; probability of passing each test $P(X_j = 1) = p$.
 - Null hypothesis and alternative hypothesis
 - $H_0: \hat{p} \geq p$; the two threads are co-located.
 - $H_1: \hat{p} < p$; the two threads are not co-located.

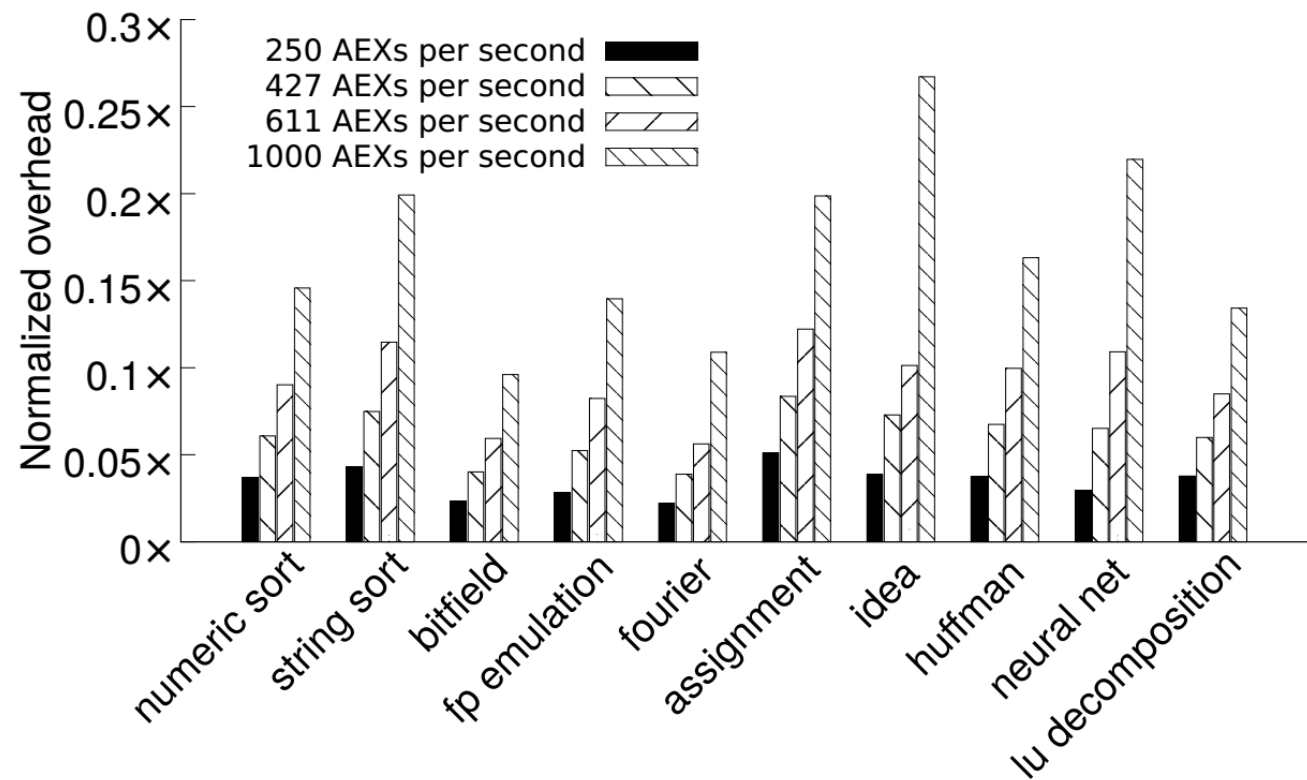
- Threat model
 - Altering the execution speed of the enclave program by (1) causing cache contention, (2) altering CPU frequency, and (3) disabling caching.

- Security analysis
 - Two threads not co-located: the co-location test will fail with very high probability
 - Two threads co-located: the co-location test will fail with very low probability

Implementation of HyperRace


- HyperRace is a compiler-assisted tool
 - A tool based on the LLVM framework
 - Instrument enclave program automatically to insert attack detection
- HyperRace working logic:
 - Create a shadow thread for each enclave thread
 - Co-location test at EENTER of the enclave thread
 - Check AEX at every basic block (of a control flow graph)
 - Reasoning: no core migration without AEX
 - May insert more checks in the same basic block if needs higher security
 - Perform co-location tests when AEXs detected

Performance overhead of co-location tests



Hardware Solutions

- Disable Hyper-Threading
 - Inform enclave developers through remote attestation
- Enable Hyper-Threading
 - Temporarily disable sibling Hyper-Thread upon EENTER, or
 - Enclave code checks co-location through MSR, and
 - Enclave code selectively disable sibling Hyper-Thread by updating MSR



Hardware-Software Co-design

- Collaboration of researchers in **architecture, system, and security**
 - Micro-architectural side channels are rooted in the computer (micro-)architecture
 - Improper assumption made by system designer for program isolation
 - Insufficient understanding of the attack models, methods, techniques
- **Key research questions & challenges:**
 - Trade-off: cost effectiveness, performance, security
 - Metrics of measurement
 - Separation of responsibilities
 - Side channels are “side” channels by definition
 - Motivate security researchers to involve in the design stage
 - Leakage-free design \neq leakage-free implementation